

# Probabilistic Programming for Scientific Discovery

Lecture 4

Ludger Paehler

*Lviv Data Science Summer School*

July 30, 2020

# Table of Contents

## Interacting with Scientific Simulators: The Engineering Challenge

Merlin

Reverb

Apache Arrow

XLA

# Outline

## Interacting with Scientific Simulators: The Engineering Challenge

Merlin

Reverb

Apache Arrow

XLA

# The Challenge

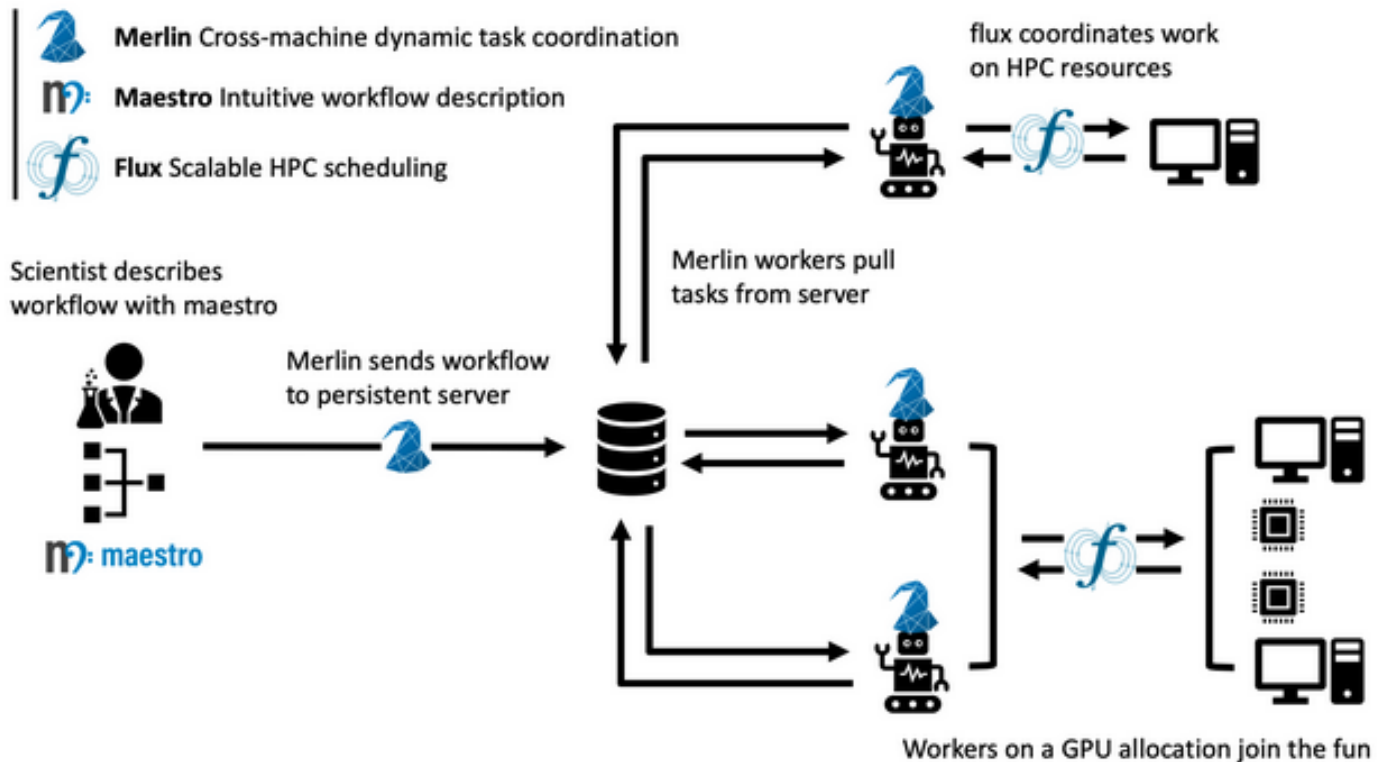
- When scaling our application/workflows to scientific applications we are faced with a multitude of problems:
  - Data movement
  - Amount of data
  - Concurrent processes, which might be prone to crashing
  - Application portability
- This is the point where probabilistic programming turns into an HPC engineering problem, where we have to choose our frameworks wisely to get the most performance out of the hardware
- Offshot: Will also run much more efficient/faster on small computers

# Merlin



- Manager for HPC-focuses simulation workflow, which can include multiple supercomputers, cloud components, as well as heterogeneous clusters.
- Made to scale to multiple hundred million simulations at a time
  - (Broken down into individual batches of simulations)
- Can automatically take advantage of spare computing capacity
- Persistent in case of individual servers failing
- Can be likened to a microservices-focussed architecture for HPC simulations

# Merlin




# Merlin


- Asynchronous task queuing library (Merlin)
- Dynamic task queuing (Celery)
- Description of our model is contained in a central workflow file
- Resilience and support for multiple users and queues (RabbitMQ)
- Scalability and fast database access (Redis)

# Merlin

`merlin run`




`merlin run-worker`





# Merlin

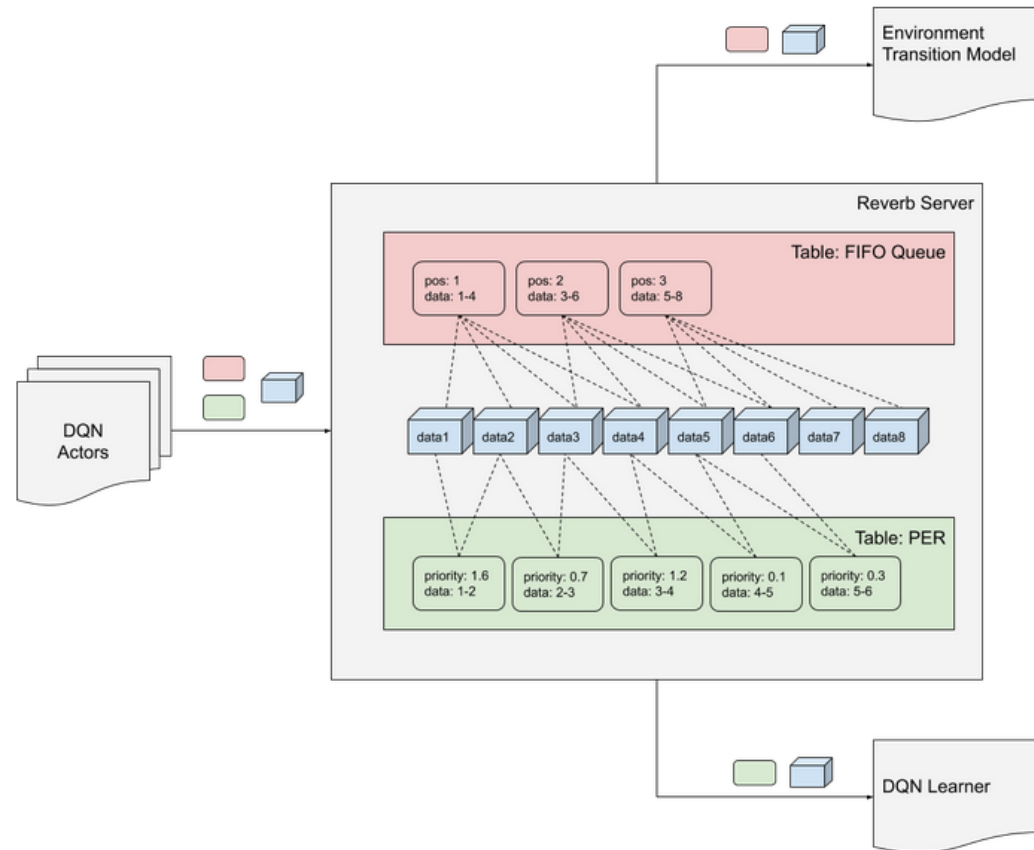


**Figure:** Any workflow we can define as a directed-acyclic graph can be performed.

# Reverb

- Experience replay system originally built for off-policy reinforcement learning
  - Can be adapted to probabilistic inference as well
- Reverb server consists of one or more tables, where tables also define sample and removal strategies, maximum item capacity, and a rate limiter.
- Multiple item items can refer to the same data element, as entries in tables only contain references to the data
- Data is deleted if there is no reference to it anymore

# Reverb



# Apache Arrow

- OSS Community initiative conceived in 2015
- Sits right at the intersection of big data, database systems, and data science tools
- Core of the project: Language agnostic open standards to accelerate in-memory computing
- Components have public APIs
- Arrow is front-end agnostic
- Dataframe backend for NVIDIA Rapids

<https://github.com/apache/arrow>


**Front end API**

**Computation Engine**

**In-memory storage**

**IO and  
Deserialization**

# Apache Arrow



# Apache Arrow

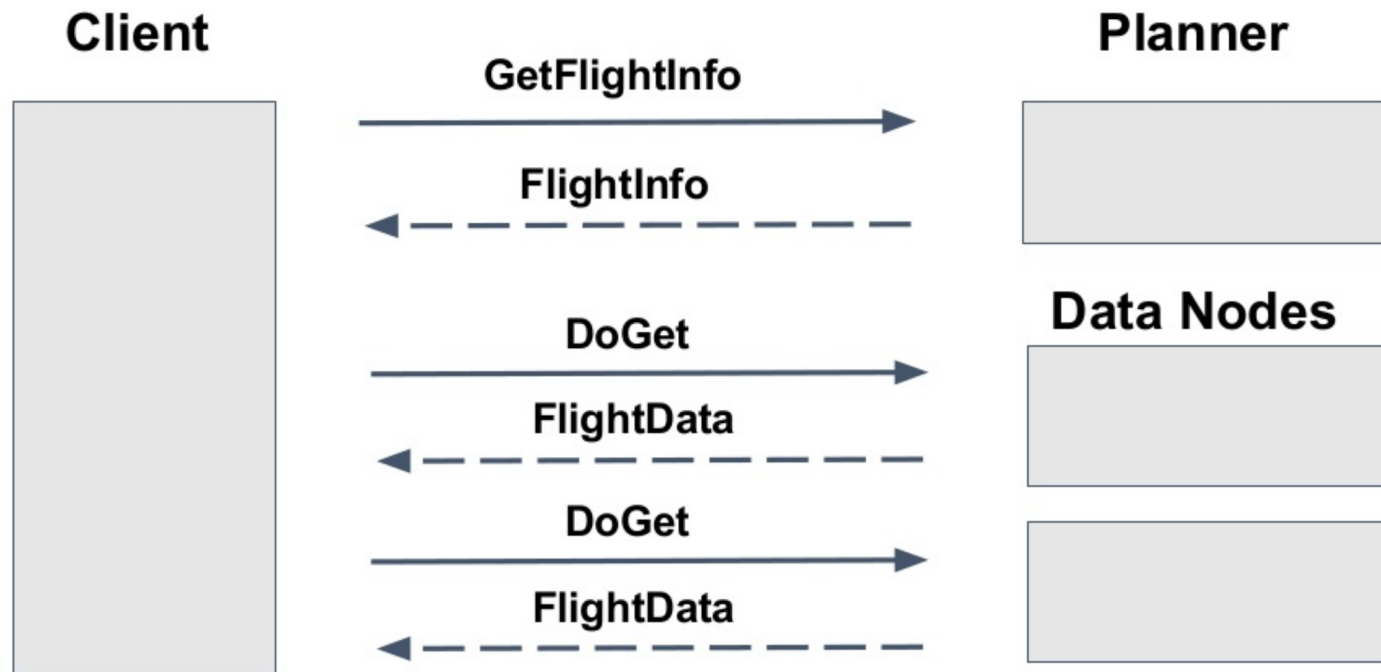
- Applied to eliminate serialization overhead in data interchange
- Improve CPU/GPU in-memory processing efficiency
- Simplify architectures
- Organized for cache-efficient access on CPUs/GPUs
- Optimized for data locality, SIMD, and parallel processing
- In-memory encoding, compression, and sparseness are in the works

# Apache Arrow

## Apache Arrow Flight

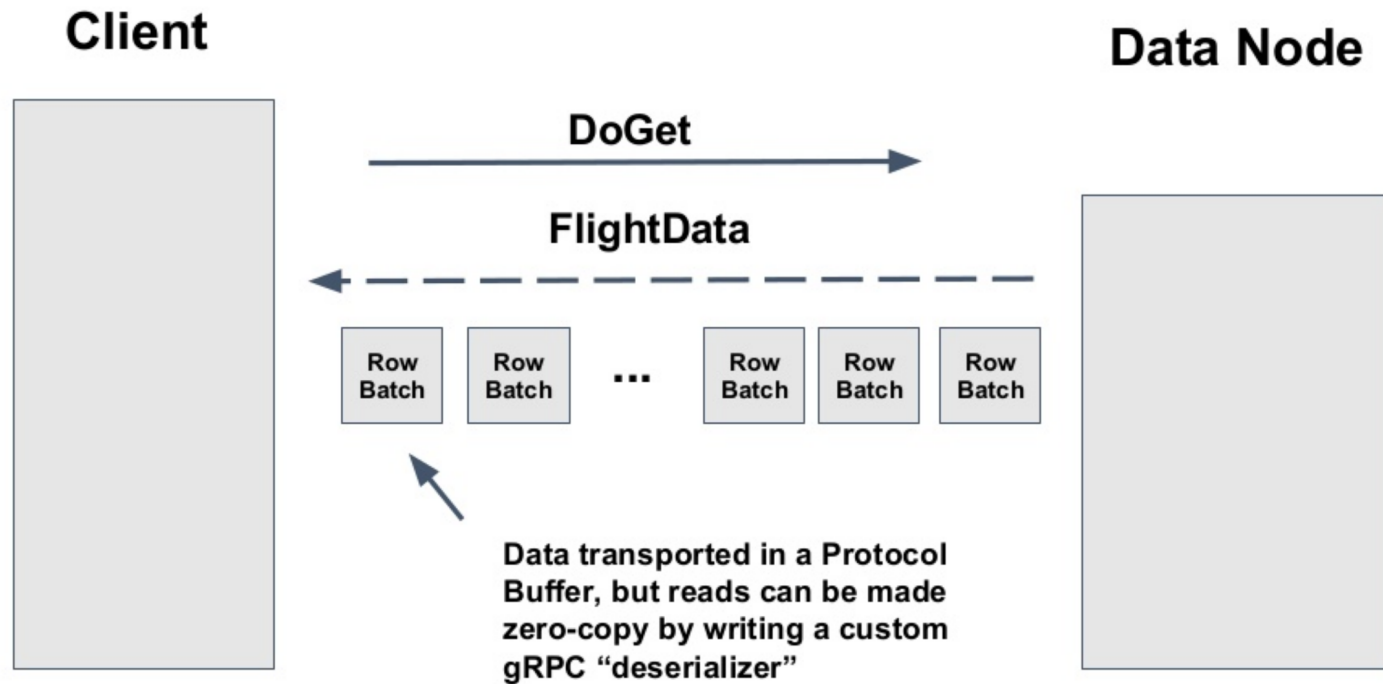
- Framework to define custom data services and receive Arrow data natively
  - Avoids any copying or deserializatio
- Low-level optimizations
- Made for continous extreme-scale datastreams, i.e. also useful for HPC applications

# Apache Arrow



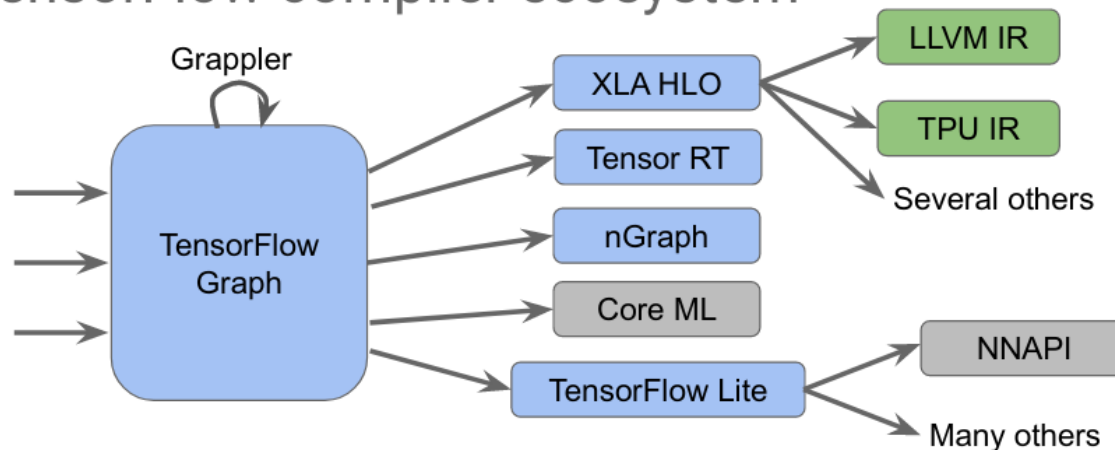


# Apache Arrow



# Machine Learning Compilers

## The TensorFlow compiler ecosystem



Many “Graph” IRs, each with challenges:

- Similar-but-different proprietary technologies: not going away anytime soon
- Fragile, poor UI when failures happen: e.g. poor/no location info, or even crashes
- Duplication of infrastructure at all levels

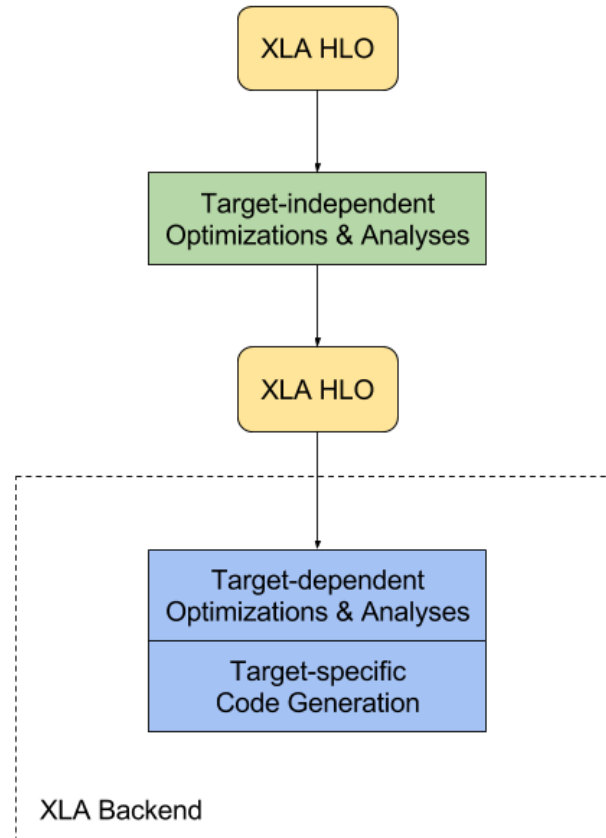


# XLA

## Purpose behind XLA:

- Improve the execution speed
- Improve memory usage
- Reduce reliance on custom operations
- Reduce mobile footprint
- Improve portability
  - Accelerator-producers just need to implement the XLA API, and our code will run seamlessly on accelerators of all types

# XLA



# The Approach

- Data movement —→ Apache Arrow
- Amount of data —→ Apache Arrow + Compression
- Concurrent processes, which might be prone to crashing —→ Merlin
- Application portability —→ XLA or MLIR